

# Heart Disease Prediction Using Neural Networks

## Case Study Report

---

### 1. Problem Description and Rationale

#### Problem Statement

Heart disease is one of the leading causes of death worldwide. Early detection can significantly improve patient outcomes through timely treatment. This project predicts whether a patient has heart disease based on 13 clinical features.

#### Why Neural Networks?

- **Non-linear relationships:** Heart disease risk factors interact in complex ways
  - **Automatic feature learning:** No need for manual feature engineering
  - **Probabilistic output:** Provides confidence scores for predictions
- 

### 2. Dataset Description and Preprocessing

#### Dataset Overview

- **Source:** Cleveland Heart Disease Dataset (UCI Repository)
- **Size:** 303 patients
- **Features:** 13 clinical attributes
- **Target:** Binary (0 = No disease, 1 = Disease)

#### Features

##### Feature Description

age      Age in years

sex      Gender (0=F, 1=M)

cp      Chest pain type (1-4)

trestbps Resting blood pressure

chol      Serum cholesterol

## Feature Description

fbs Fasting blood sugar > 120

restecg Resting ECG results

thalach Max heart rate achieved

exang Exercise induced angina

oldpeak ST depression

slope Slope of ST segment

ca Number of vessels colored

thal Thalassemia type

## Preprocessing Steps

### 1. Normalization

```
Model.normalizeMinMax(X);
```

Applied Min-Max scaling to normalize all features to [0, 1] range. This ensures:

- All features contribute equally to learning
- Faster convergence during training
- Better performance with activation functions

### 2. Train-Test Split

```
SplitResult split = Model.trainTestSplit(X, Y, 0.3);
```

- Training: 70% (212 samples)
- Testing: 30% (91 samples)

---

### 3. Neural Network Architecture

#### Network Structure

**Architecture: 13 → 32 → 16 → 1**

```
model.createNetwork(3,
```

```
    new int[] { 13, 32, 16 },
```

```
new int[] { 32, 16, 1 });
```

Input Layer: 13 neurons (matches 13 features)

Hidden Layer 1: 32 neurons + ReLU

Hidden Layer 2: 16 neurons + ReLU

Output Layer: 1 neuron + Sigmoid

## Why This Architecture?

### Layer Sizes:

- **13 → 32**: First layer expands to capture feature combinations and interactions
- **32 → 16**: Second layer compresses information, learning important patterns
- **16 → 1**: Final layer produces single probability output

This creates a bottleneck structure that forces the network to learn the most important features.

### Activation Functions

#### Hidden Layers: ReLU

```
model.setActivationFunction(new ReLU());
```

#### Why ReLU for hidden layers?

- **Fast training**: Simple computation ( $\max(0, x)$ )
- **No vanishing gradients**: Gradient is either 0 or 1
- **Sparse activation**: Some neurons turn off, making the network efficient
- **Standard choice**: Works well for most classification tasks

#### Output Layer: Sigmoid

```
model.getNetwork().getLayers()[2].setActivationFunction(new Sigmoid());
```

#### Why Sigmoid for output?

- **Probability output**: Produces values between 0 and 1
  - **Binary classification**: Perfect for yes/no decisions (disease or no disease)
  - **Interpretable**: Output can be read as "confidence" percentage
-

## 4. Hyperparameter Configuration

```
model.setWeightInitializer(new XavierNormal());  
model.setLossFunction(new BinaryCrossEntropy());  
model.setActivationFunction(new ReLU());  
model.setLearningRate(0.003);  
model.setBatchSize(32);  
model.setEpoches(5000);
```

Parameter	Value	Why?
Learning Rate	0.003	Balanced speed - not too fast (unstable) or slow (takes forever)
Batch Size	32	Gives ~7 batches per epoch for smooth learning
Epochs	5000	Enough iterations to fully converge
Weight Init	Xavier Normal	Keeps activations stable across layers
Loss Function	Binary Cross-Entropy	Standard for binary classification, better gradients than MSE

### Why Xavier Initialization?

Initializes weights based on layer sizes:  $W \sim N(0, \sqrt{2/(input\_size + output\_size)})$

- Prevents vanishing/exploding gradients
- Helps network learn from the start

### Why Binary Cross-Entropy Loss?

Formula:  $L = -[y * \log(\hat{y}) + (1-y) * \log(1-\hat{y})]$

- Strongly penalizes wrong confident predictions
- Perfect match for Sigmoid output (both deal with probabilities)
- Better than MSE for classification

---

## 5. Training and Results

### Training Process

- **Training samples:** 212 patients
- **Test samples:** 91 patients
- **Total epochs:** 5000
- **Gradient updates per epoch:** ~7 (212/32)

### **Performance Metrics (from one of the trained models)**

Metric	Value
<b>Accuracy (total correct)</b>	90.00%
<b>Precision (correct when predicting disease)</b>	88.10%
<b>Recall (% of actual disease cases caught)</b>	90.24%

### **Confusion Matrix**

	Predicted No	Predicted Yes
Actual No	40	5
Actual Yes	4	42

### **What this means:**

- **40 True Negatives:** Correctly identified healthy patients
- **42 True Positives:** Correctly identified disease patients
- **5 False Positives:** Healthy patients flagged as sick (safe error)
- **4 False Negatives:** Disease patients missed (critical error)

**Note:** we made that predictions that are equal or higher than 0.5 are equal to 1 or equal to that person is having a heart disease and predictions that are less than 0.5 are equal to 0 or equal to that person is not having a heart disease


### **Loss Curve**

**-Rapid learning\*\*:** Loss drops sharply from 0.72 to 0.40 in first 500 epochs

**-Smooth convergence\*\*:** Gradual decrease to ~0.25 with no oscillations

**\*\*Stable training\*\*:** No signs of overfitting or instability

**\*\*Optimal stopping\*\*:** Plateau around epoch 3000 indicates sufficient training



## 6. How the Library Was Used

### Step 1: Load and Preprocess Data

(**CSVHelper is not part of the library**)

```
double[][] X = CSVHelper.selectColumns("Heart_disease_cleveland_new.csv",
    new int[] { 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12 }, 1);
double[][] Y = CSVHelper.selectColumns("Heart_disease_cleveland_new.csv",
    new int[] { 13 }, 1);
```

(the number besides the column numbers is how many rows we want to skip its useful for skipping feature names)

```
Model.normalizeMinMax(X);
SplitResult split = Model.trainTestSplit(X, Y, 0.3);
```

## **Step 2: Configure Model**

```
Model model = new Model();
model.setWeightInitializer(new XavierNormal());
model.setLossFunction(new BinaryCrossEntropy());
model.setActivationFunction(new ReLU());
model.setLearningRate(0.003);
model.setBatchSize(32);
model.setEpoches(5000);
```

## **Step 3: Build Network**

```
model.createNetwork(3,
    new int[] { 13, 32, 16 },
    new int[] { 32, 16, 1 });
```

```
model.getNetwork().getLayers()[2].setActivationFunction(new Sigmoid());
```

## **Step 4: Train**

```
model.train(split.X_train, split.Y_train);
```

## **Step 5: Evaluate**

```
double[][] predictions = model.predict(split.X_test);
```

```
// Calculate metrics
```

```
for (int i = 0; i < predictions.length; i++) {
    int pred = predictions[i][0] >= 0.5 ? 1 : 0;
    int actual = (int) split.Y_test[i][0];
    // Count TP, TN, FP, FN...
}
```

## **Step 6: Save Model and test data**

```
CSVHelper.writeCSV("X_test.csv", split.X_test);
```

```
CSVHelper.writeCSV("Y_test.csv", split.Y_test);  
model.saveModel("model.nn");
```


If we want to use the model directly without retraining it we can just use Model  
model = Model.loadModel("model.nn"); and load the test data back ( double[][] X\_test =  
CSVHelper.selectColumns("X\_test.csv",

```
    new int[] { 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12 }, 0);  
    double[][] Y_test = CSVHelper.selectColumns("Y_test.csv", new int[] { 0 }, 0);  
)  
(we didn't skip the feature names this time because it does not exist in the test data csv  
file)
```

---

## 7. Results Screenshots

### Loss curve graph picture



### Results console output:

```
Accuracy (total correct): 90.0%
Precision (correct when predicting disease): 88.09523809523809%
Recall (% of actual disease cases caught): 90.2439024390244%
```